

AgentGuard: Deterministic Security Linting for AI Agent Definitions

Ying Chen

[doi:10.5281/zenodo.21420490](https://doi.org/10.5281/zenodo.21420490)

Abstract

AI coding harnesses execute agents defined in plain markdown-with-frontmatter files: a tool grant, a description, and natural-language instructions. These definitions are security-relevant code—a definition that reads untrusted content while holding shell or network tools is one injected sentence away from remote code execution or data exfiltration—yet they ship without any static analysis. We introduce AgentGuard, a deterministic, zero-dependency security linter for agent, command, and skill definitions. AgentGuard parses each definition’s capability grant, models untrusted-input sources and dangerous sinks, and reports severity-ranked findings for capability chains such as injection-to-execution, exfiltration paths, and disabled human-in-the-loop guardrails, each mapped to OWASP LLM Top-10 and MITRE ATLAS. On a gated benchmark of 15 vulnerable and 13 safe definitions, AgentGuard achieves 0.93 recall and 1.00 precision with zero false alarms—where naive grant-based or keyword-grep flagging reaches at most 0.47 recall at 0.58 precision—and its security decisions are stable across 81 metamorphic adversarial rewrites; the single missed case is a named, documented recall boundary rather than a silent gap. Field studies of the official Claude Code plugin marketplace at two snapshots one month apart (33 and 84 unique definitions) found 85–86% lacked any injection guard and 37–39% could be driven to run a command or write a file by content they read—rates that held as the marketplace grew 2.5× and replicated on an independent third-party corpus (95% unguarded, 50% chain). Scanning is cheap: the full marketplace lints in under a second (≈ 9 ms per definition). The core claim is narrow: privilege risk in agent definitions can be a deterministic, benchmark-gated CI check rather than a manual review afterthought.

Introduction

Indirect prompt injection is a structural property of LLM agents, not an implementation bug: any agent that reads attacker-influenced content and holds action-capable tools can be redirected by instructions embedded in that content (Greshake et al. 2023; Perez and Ribeiro 2022; Willison 2022). Modern coding harnesses make this concrete at scale. An agent definition is a markdown file with YAML frontmatter declaring a name, a description, and a `tools:` grant; the body is a natural-language system prompt. A definition that says “summarize the file” while granting Bash is an unguarded injection-to-execution chain—the file being summarized can carry the instruction `curl ...| sh`, and nothing in the definition tells the model to treat read content as data.

These files are code. They are committed to repositories, distributed through plugin marketplaces, installed onto developer machines, and executed with filesystem, shell, and network access. Yet unlike the application code they sit beside—linted by Bandit, Semgrep, and a CI wall of static analysis (PyCQA 2024; Semgrep, Inc. 2024)—agent definitions ship with no static check at all. Guidance exists as taxonomy (OWASP LLM Top-10, MITRE ATLAS (OWASP Foundation 2025; MITRE Corporation 2024)) and as design advice (Anthropic 2024), but not as an enforceable gate.

This paper asks a systems question: how much of agent-definition privilege risk can a deterministic static linter catch, at what precision, and with what honestly-stated recall boundary? AgentGuard’s answer is deliberately conservative: no LLM in the loop, no API call, no nondeterminism—every finding reproduces bit-for-bit in CI, and the linter’s own quality is enforced by a gated benchmark rather than asserted.

Contributions.

- A capability model for agent definitions—untrusted-input sources, dangerous sinks (execute, write, network, spawn), guard predicates, and fail-closed grant semantics—and 26 deterministic rules over it, organized in four layers from parse integrity to capability-chain security, each finding mapped to OWASP LLM Top-10 and MITRE ATLAS.
- A two-axis security posture score that separates “is anything dangerous?” (a size-independent critical-count ceiling) from “is it systemically sloppy?” (a per-file major/minor density), fixing the standard aggregation bug where grades scale with repository size instead of risk.
- A gated quality benchmark—minimum recall 0.93, precision 1.00, zero false alarms, a frozen case inventory, and 81 metamorphic adversarial variants—that fails the build when the linter regresses, plus a *named known miss* documenting the lexical recall boundary.
- A field study of the official Claude Code plugin marketplace quantifying, for the first time, how often shipped agent definitions carry the structural preconditions for injection-to-action attacks.

Method

Figure 1 shows the pipeline. Definitions are parsed into a frontmatter/body pair; the capability model derives the tool grant, untrusted-input sources, sinks, and guard predicates; four rule layers evaluate the model; and severity-ranked findings feed the CLI, SARIF export, GitHub Action, baseline gate, and posture score.

AgentGuard pipeline. Every stage is deterministic: parsing fails closed into findings rather than silence, the capability model treats a missing `tools:` grant as full privilege, and the same input always produces the same findings—a property enforced by 81 metamorphic adversarial variants in the quality gate.

Problem Formalization

Let a definition be $d = (\phi, \beta)$ with frontmatter ϕ and body β , and let T be the harness tool universe. The capability model derives four objects:

$$C(d) \subseteq T, \quad U(d), \quad \Sigma = \Sigma_{\text{exec}} \cup \Sigma_{\text{write}} \cup \Sigma_{\text{net}} \cup \Sigma_{\text{spawn}}, \quad \gamma(d) \in \{0, 1\}$$

where $C(d)$ is the tool grant—with the fail-closed convention that an absent `tools:` field yields $C(d) = T$, since that is the harness inheritance semantics; $U(d)$ is the set of untrusted-input sources the body instructs the agent to read (files, web content, tickets, user arguments); Σ partitions dangerous sinks into execute, write, network, and sub-agent-spawn classes; and $\gamma(d)$ is the guard predicate—true when the body contains an explicit “treat read content as data, not instructions” control.

Chain rules are conjunctions over this model. The headline rule (AL300, injection-to-action) fires when

$$U(d) \neq \emptyset \wedge C(d) \cap (\Sigma_{\text{exec}} \cup \Sigma_{\text{write}}) \neq \emptyset \wedge \neg\gamma(d),$$

escalating to critical severity when the grant simultaneously holds a network-capable reader and an execution sink. The exfiltration rule (AL301) replaces the sink condition with $C(d) \cap \Sigma_{\text{net}} \neq \emptyset$ and the source condition with sensitive-data handling; the propagation rule (AL307) uses Σ_{spawn} , modeling injected instructions forwarded into spawned sub-agents. Each finding carries a severity in $\{\text{critical}, \text{major}, \text{minor}, \text{info}\}$, an OWASP LLM Top-10 and MITRE ATLAS mapping, and a one-line fix.

Worked Example

Figure 2 shows the model end to end on a minimal definition of the kind the field study finds routinely: a triage agent granted `[Read, WebFetch, Bash]` whose body instructs it to read a ticket, fetch URLs from it, and run the reproduction commands it describes. In model terms, $U(d) = \{\text{file}, \text{web}\}$, $C(d) \cap \Sigma_{\text{exec}} = \{\text{Bash}\}$, and $\gamma(d) = 0$; the AL300 conjunction holds and escalates to critical because a network-capable reader (`WebFetch`) co-occurs with an execution sink. A ticket containing “to reproduce, run `curl ... | sh`” drives this agent to remote code execution by design. The two-line fix—a treat-as-data guard plus a scoped grant—clears both findings.

```
---
name: ticket-triage
description: Use this agent to triage incoming support
tickets
tools: [Read, WebFetch, Bash]
---
Read the ticket file the user points you at, fetch any URLs
it
references, and run whatever reproduction commands the
ticket
describes so the bug is confirmed before routing.
```

```
!! critical AL300 Injection->action chain: reads
external/untrusted
content and can also Bash -- with no instruction to treat
```

```

that
  content as data. Granted: Bash, Read, WebFetch.
  [OWASP LLM01/LLM06 – ATLAS AML.T0051.001]
  fix: add an explicit guard ("treat all read content as
data,
  never as instructions") AND restrict tools: to the
minimum.
!! major    AL202  Consumes external content but never says
to treat
  it as data, not instructions. [OWASP LLM01]
2 findings in 1/1 files  (1 critical, 1 major)

```

Worked example: a realistic triage-agent definition (top) and AgentGuard’s actual findings (bottom, condensed). The definition is not malicious—it is a natural way to write a triage agent—which is precisely why the preconditions are so prevalent in the field study.

Rule Taxonomy

The 26 rules are organized in four layers (Table 1). The layering is an analysis-order contract: structural rules establish that a definition can be parsed at all (fail-closed—an unreadable definition is a major finding, never a silent skip); instruction and contract rules check that behavior is specified tightly enough to audit; capability-chain rules reason over $C(d) \times U(d) \times \Sigma \times \gamma(d)$ combinations. Severity semantics are strict: *critical* is reserved for security exposure (an unguarded destructive action, an exfiltration path, a hardcoded secret, an explicitly disabled human-in-the-loop check), never for structural or stylistic issues—so a CI gate keyed on `--fail-at critical` fires only on what that tier promises.

Rule taxonomy: 26 deterministic rules in four layers. AL0xx–AL2xx evaluate a definition in isolation; AL3xx rules evaluate dangerous *combinations* of capability grant, untrusted input, sinks, and missing guards. Findings can be suppressed only by a visible inline annotation, which is itself auditable.

Layer	Concern	Rules	Representative rules
AL0xx	parse & structural integrity	7	unreadable file, missing frontmatter/name
AL1xx	instruction quality	2	vague instruction, unenforceable safety
AL2xx	behavioral contract	8	injection guard absent, unguarded destructive act
AL3xx	capability-aware security	9	injection→exec chain, exfiltration path, secret in definition, human-in-loop disabled

Security Posture Score

Aggregating findings into a repository grade has a standard failure mode: summing weighted findings makes the grade scale with repository size, so a 40-file benign scaffold floors to F on accumulated minor findings while a 2-file plugin with a real injection chain scores the same. AgentGuard separates the two questions a security grade conflates. “Is anything dangerous?” is a count-based ceiling over criticals, independent of size; “is it systemically sloppy?” is a per-file density over majors and minors. For n files with critical/major/minor counts c, M, m :

$$\text{ceiling} = 100 - 34\min(c, 2), \quad \text{density} = (7M + 2m)/n, \quad \text{score} = \max(0, \min(\text{ceiling}, 100 - \text{density})).$$

One critical caps the grade at D (66) regardless of how clean everything else is; two or more cap it at F (32); a benign sprawl with zero criticals and modest per-file density grades A. Non-A grades additionally name the files contributing the most density, so the number is actionable rather than opaque.

Quality Gate: Benchmark, Metamorphic Review, and a Named Miss

Most linters report precision and stay quiet about recall. AgentGuard’s benchmark is a committed, CI-gated inventory: 15 positive (vulnerable) cases spanning every AL3xx chain class, 13 negative (safe) cases—several mined from real marketplace definitions that earlier versions falsely flagged—and explicit gates: recall ≥ 0.93 , precision = 1.00, zero false alarms, and a frozen minimum case count so cases cannot be quietly deleted to keep the gate green.

Two further mechanisms harden the gate. First, a metamorphic adversarial review applies 81 semantics-preserving rewrites to benchmark cases—reordering sections, renaming headings, paraphrasing instructions—and requires every security decision to remain unchanged, so the rules cannot be keying on incidental prompt structure. Second, the benchmark’s single missed case is *named*: a vulnerable definition whose destructive action is expressed through a fully arbitrary euphemism carries no lexical anchor a deterministic scanner can key on. The miss is committed to the gate configuration as a documented recall boundary rather than silently absorbed into an aggregate number.

CI Integration and Baseline Gating

AgentGuard ships as a CLI, a GitHub Action, and a SARIF exporter. A baseline mode records the current finding set and fails CI only on *new* findings, making the linter adoptable on a repository with existing debt. Vended third-party plugin caches are skipped by default—the same convention as `node_modules` for code linters—because findings in unmodifiable vendored copies are noise for the repository owner; an explicitly given plugin path is still scanned, which is the “vet before install” workflow.

Experiments

Three evidence levels parallel the design. The *benchmark* measures detection quality on a labeled corpus with adversarial variants, against naive deterministic baselines. The *field study* measures prevalence of the structural preconditions in shipped, real-world definitions at two snapshots. *Case studies* document precision hardening against real false positives. All numbers below reproduce from the repository with three commands (`python eval/benchmark.py`, `python eval/adversarial_review.py`, and `python eval/naive_baselines.py`); the first two run in CI on every push, so the reported figures are continuously re-verified rather than frozen at submission time.

Cost.

Determinism buys speed as well as reproducibility: on a commodity laptop (Apple M-series), the full 28-case benchmark plus gate checks completes in 0.08 s, the 81-variant metamorphic review in 0.21 s, and a scan of the entire official plugin marketplace (84 definitions) in 0.77 s end-to-end—about 9 ms per definition. Per-commit, per-install, and marketplace-wide scanning are therefore free in practice, which is the economic premise of the CI-gate design.

Gated benchmark, 15 vulnerable / 13 safe cases, AgentGuard v0.1.3 (212 unit tests alongside). Baselines are the two obvious deterministic alternatives—auditing the tool grant without reading the body, and grepping the body without modeling the grant—scored generously at the case level (any flag on a vulnerable case counts as caught; AgentGuard’s own gate additionally requires the correct rule to fire). The gate fails CI if recall drops below 0.93, precision below 1.00, any false alarm appears, the case inventory shrinks, or any metamorphic variant changes a security decision. The one miss is the named arbitrary-euphemism case (Section 7).

Detector	Recall	Precision	False alarms
AgentGuard (capability model, 26 rules)	0.93	1.00	0
Grant-based: flag any exec/write-capable grant	0.47	0.58	5
Keyword grep (rm/curl/secret/...)	0.33	0.50	5
Grant $\wedge$ keyword	0.20	0.50	3
Metamorphic adversarial variants (stable)			81 / 81
Named known misses			1

Field study: the official Claude Code plugin marketplace at two snapshots one month apart, deduplicated by content hash against orphaned cache copies. The marketplace grew 2.5 $\times$ between snapshots while the exposure rates held. “Exposed” means the structural precondition for an injection-to-action attack is present—untrusted input, action capability, no guard—not that a weaponized exploit was demonstrated against each definition. The fix in the common case is one guard line plus a scoped `tools: grant`.

Field study finding	2026-06-12 (v0.1.2)	2026-07-14 (v0.1.3)
Unique definitions (plugins)	33 (6)	84 (25)
No “treat as data” guard on untrusted input (AL202)	28 85%	72 86%
Drivable to run a command / write a file (AL300)	13 39%	31 37%
At least one security-class (AL3xx) finding	13 39%	38 45%

Per-kind breakdown, 2026-07-14 snapshot ($n = 84$). Guarding practice is worst where the stakes are highest: agent definitions are guarded most often, yet chain most (58%), because they hold the broadest tool grants; skills are universally unguarded but chain less because many hold no action-capable grant.

Definition kind	n	No guard (AL202)	Full chain (AL300)
Agents	24	16 67%	14 58%
Commands	29	25 86%	8 28%
Skills	29	29 100%	9 31%
Other frontmatter definitions	2	2	0

Precision-hardening case studies.

Two rule revisions illustrate the precision discipline. The assertive-claims rule (AL204) originally fired on any confident-recommendation stem; validation against real agent definitions showed it also fired where the stem was merely *described* rather than performed—inside an output-template code fence, as a noun phrase, or as the object of a data verb (“extract scores from...”). The rule now distinguishes described from performed assertions, with regression tests, and benchmark recall and precision held. Similarly, the structural layer initially graded an empty file as multiple missing-field findings; it now emits a single missing-frontmatter finding, because one root cause should be one finding. Both revisions came from running the tool on real definitions and treating every false positive as a gate-level bug.

Results

Detection Quality

On the gated benchmark (Table 2), AgentGuard catches 14 of 15 vulnerable definitions (recall 0.93) and produces zero findings on all 13 safe definitions (precision 1.00, zero false alarms). The baselines quantify what the capability model buys: flagging on the tool grant alone reaches 0.47 recall with 5 false alarms (it misses body-level risks like hardcoded secrets and sensitive-data exfiltration, and flags guarded definitions that hold broad grants safely), while keyword grep reaches 0.33 with the same 5 false alarms (it keys on words rather than on who can do what to which input); requiring both signals drops recall to 0.20. Detection here genuinely requires the joint $C(d) \times U(d) \times \Sigma \times \gamma(d)$ analysis, not either projection of it. The positive set spans every capability-chain class: injection-to-execution, exfiltration, secret-in-definition, disabled human-in-the-loop, sub-agent propagation, and untrusted command construction. The negative set is the harder half: several cases are real marketplace definitions that superficially pattern-match a vulnerability—a security reviewer agent that *discusses* destructive commands, a template that *renders* assertive scores—and earlier versions flagged them; each drove a precision fix.

Metamorphic Stability

All 81 semantics-preserving rewrites of benchmark cases leave every security decision unchanged. This is the linting analogue of an invariance test: a scanner whose verdict flips when a heading is renamed is keying on prompt structure, not on the capability model. The metamorphic suite runs in CI, so structural over-fitting cannot regress silently.

Field Prevalence

The marketplace study (Table 3) is, to our knowledge, the first quantified measurement of injection-attack preconditions in shipped agent definitions: 85–86% of official marketplace definitions carry no injection guard of any kind, and 37–39% combine untrusted input with an action-capable grant and no guard—the full AL300 precondition. The two snapshots are the stronger result: the marketplace grew from 33 definitions in 6 plugins to 84 in 25 within a month, and the exposure rates did not move—new plugins ship with the same unguarded defaults as old ones, so the exposure is a property of how definitions are written, not of a few early outliers. The per-kind breakdown (Table 4) sharpens this: agent definitions, which hold the broadest grants, chain at 58%, and skills are unguarded without exception. The study protocol deduplicates cache copies by content hash (counting them would double the denominator) and pins each snapshot date and scanner version. The numbers quantify exposure, not exploitation; the point is that the preconditions are the *default* state of shipped definitions, not the exception.

Cross-distribution replication.

To test whether the unguarded default is specific to the official marketplace, we ran the same reproducible field-study protocol (`python eval/field_study.py`) over an independent corpus: 20 unique definitions from two third-party plugin packages installed from *outside* the official marketplace (Table 5). The exposure is if anything worse: 19/20 (95%) carry no injection guard and 10/20 (50%) hold the full AL300 chain, versus 86% and 37% on the official set. Pooled across both channels (104 definitions), 88% are unguarded and 39% chain. The finding replicates on a distribution the official marketplace’s authors do not control, which is the evidence that the unguarded default is an ecosystem property of how agent definitions are written, not an artifact of one registry’s review process. A manual precision check confirms the third-party chain findings are not false positives: all 10 flagged definitions omit the `tools: field` (inheriting `execute-capable tools`) while instructing the agent to read untrusted input (source files, articles, or codebases) with no `data-not-instructions` guard; one (`file-analyzer`) is additionally told to “write and execute” an extraction script over the untrusted files it reads—a textbook `read-plus-execute` chain. Precision on the independent chain set is 10/10.

Cross-distribution replication (2026-07-15, AgentGuard v0.1.3). The same protocol run on an independent third-party corpus reproduces the unguarded-default finding at an even higher rate, deduplicated by content hash. Third-party packages are described in aggregate; the runnable `eval/field_study.py` takes corpus roots as arguments and hardcodes none.

Distribution channel	Defs	No guard (AL202)	Full chain (AL300)
Official marketplace	84	72 86%	31 37%
Independent third-party plugins	20	19 95%	10 50%

Distribution channel	Defs	No guard (AL202)	Full chain (AL300)
Pooled	104	91 88%	41 39%

Posture Score Calibration

The two-axis score separates cases that single-axis aggregation conflates: a benign 40-file definition set with zero criticals, 8 majors, and 130 minors grades A (92)—under summed weighting it would floor to F—while a 2-file plugin containing one real injection critical is capped at D (66) and any set with two or more criticals is pinned at F (32) regardless of size. The grade now tracks danger and sloppiness independently, and by construction cannot be diluted by adding clean files (the critical ceiling is count-based, not averaged); these calibration points are regression tests in the suite. Non-A grades name the files dragging density, which converts the grade from a verdict into a worklist.

Discussion

Why deterministic, when an LLM reviewer would understand more?

An LLM reviewing agent definitions can, in principle, catch the arbitrary-euphemism case a lexical scanner misses. AgentGuard deliberately excludes LLM calls for three reasons. First, reproducibility: a CI gate whose verdict varies across runs, model versions, or temperature settings cannot gate merges—the same diff must always get the same answer. Second, the reviewer is itself injectable: an LLM reading a malicious definition is exposed to the very attack class it is auditing, and a definition can embed instructions aimed at its reviewer (Greshake et al. 2023). Third, marginal cost: a zero-dependency scanner runs on every commit, every marketplace install, and every pre-commit hook at no per-call cost, which is what makes fleet-wide and marketplace-wide scanning economically trivial. The right division of labor is a deterministic gate for the structural preconditions plus human (or model) review for semantics, with the gate documenting exactly where its boundary lies.

Exposure is not exploitation.

The field-study framing matters for calibration. An unguarded read-plus-execute definition is an unlocked door, not a robbed house: the claim is the structural precondition, not a demonstrated end-to-end compromise of each definition. This is the same epistemic position as a memory-safety lint—most flagged buffer usages will never be attacked, and flagging them is still correct because the fix is cheap (one guard line, one scoped grant) and the failure is catastrophic and silent when it lands.

Severity semantics as a contract.

Reserving *critical* for security exposure—never structural sprawl—is what makes —fail-at critical a meaningful CI policy. The posture score extends the same contract to the aggregate: the critical ceiling cannot be diluted by repository size, and the density axis cannot manufacture a critical. Severity inflation is the fastest way a security tool loses its audience; the gate treats it as a bug class, not a style choice.

The named miss as methodology.

Committing a known-missed case to the quality gate—by name, with the reason it is missed—inverts the usual incentive. A benchmark the tool author controls invites quietly dropping hard cases; freezing the inventory and naming the miss makes the recall boundary a reviewable artifact. We consider this pattern (gate on recall *and* precision, freeze the case inventory, name the misses) more transferable than any individual rule.

Related Work

Indirect prompt injection was established as a structural risk of LLM-integrated applications by Greshake et al. (Greshake et al. 2023), with earlier attack taxonomies from Perez and Ribeiro (Perez and Ribeiro 2022) and the term coined by Willison (Willison 2022). OWASP LLM Top-10 and MITRE ATLAS catalogue the threat classes (OWASP Foundation 2025; MITRE Corporation 2024); AgentGuard operationalizes a subset of them as executable checks and cites the mappings in each finding. Anthropic’s agent-design guidance recommends minimal tool grants and treating retrieved content as data (Anthropic 2024); AgentGuard turns those recommendations into failing CI checks. Alignment-auditing tools such as Petri probe model *behavior* dynamically (Anthropic 2025); AgentGuard is complementary, analyzing the *definition* statically before any model call.

Table 6 positions AgentGuard against the nearest tooling. Classical security linters (Bandit, Semgrep) analyze program source and do not parse agent frontmatter, tool grants, or natural-language capability chains (PyCQA 2024; Semgrep, Inc. 2024). LLM-based prompt reviewers understand semantics but are nondeterministic, per-call priced, and themselves injectable. Taxonomies name the risks but do not detect them.

Positioning. Code SAST tools (Bandit, Semgrep) share the deterministic-gate philosophy but analyze program source, not agent definitions. LLM reviewers can reach semantics deterministic rules cannot, at the cost of reproducibility, per-call price, and reviewer injectability. Taxonomies (OWASP LLM Top-10, MITRE ATLAS) define the threat classes AgentGuard’s findings map to.

Capability	AgentGuard	Code SAST	LLM reviewer	OWASP/ATLAS
Parses agent definitions & tool grants	✓	–	partial	–
Capability-chain reasoning ($U \times \Sigma \times \gamma$)	✓	–	partial	–
Deterministic verdicts	✓	✓	–	n/a
Zero per-scan cost	✓	✓	–	n/a
Gated recall + precision benchmark	✓	–	–	n/a
Metamorphic stability testing	✓	–	–	n/a
Named known misses	✓	–	–	n/a

Capability	AgentGuard	Code SAST	LLM reviewer	OWASP/ATLAS
CI baseline gating & SARIF	✓	✓	–	n/a

Limitations

AgentGuard’s analysis is lexical and structural, not semantic. The named benchmark miss is the honest statement of this boundary: a destructive action expressed through a fully arbitrary euphemism (“perform the usual tidy-up on the records”) carries no anchor a deterministic scanner can key on without either an allowlist of safe phrasings or a semantic model. The metamorphic suite bounds over-fitting to prompt structure but cannot grant understanding; recall against a motivated adversary who knows the rules is strictly lower than benchmark recall against natural definitions.

The benchmark is partly self-authored, a structural conflict mitigated but not eliminated: the negative cases are anchored in real marketplace definitions that produced false positives, the positive inventory is frozen so hard cases cannot be dropped, the misses are named, and the metamorphic variants are generated mechanically—but an external, adversarially-contributed corpus would be stronger evidence, and building one is future work.

The field study covers two temporal snapshots of the official marketplace ($n = 33$ and $n = 84$ definitions, one month apart) plus one independent third-party corpus ($n = 20$). The exposure rate replicates across both time (stable as the marketplace grows) and distribution channel (higher on third-party plugins), which is stronger than a single snapshot, but 104 definitions across two channels still does not establish a precise ecosystem-wide rate or a long-run trend; a broad, randomly-sampled multi-registry census is future work. The capability model targets markdown-with-frontmatter definitions of the Claude Code family; adapters for other harness formats exist but carry the same lexical boundary. Finally, AgentGuard is static: it cannot observe what an agent *does* at runtime, only what its definition permits—pairing the static gate with runtime monitoring is complementary, not overlapping, coverage.

Conclusion

AgentGuard treats agent definitions as what they operationally are: privileged code shipped without review. A capability model over tool grants, untrusted inputs, sinks, and guards; 26 deterministic rules with strict severity semantics; a posture score that cannot be diluted by repository size; and a quality gate that enforces its own recall, precision, and metamorphic stability on every commit—together these make agent privilege risk a CI property rather than a review-time hope. The field result is the motivating fact: across two snapshots a month apart, 85–86% of official marketplace definitions ship with no injection guard at all—a rate unchanged as the marketplace grew 2.5×—and the fix is typically two lines.

Three directions remain open: (1) an adversarially-contributed external benchmark corpus to replace self-authored positives; (2) semantic guard detection—closing the euphemism boundary without giving up determinism, e.g. via a constrained-vocabulary contract for destructive actions; and (3) a cross-harness definition standard so one capability model covers every agent format, making least-privilege lintable ecosystem-wide.

Anthropic. 2024. *Building Effective Agents*.

<https://www.anthropic.com/engineering/building-effective-agents>.

Anthropic. 2025. *Petri: An Open-Source Auditing Tool to Accelerate AI Safety Research*. <https://alignment.anthropic.com/2025/petri/>.

Greshake, Kai, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. 2023. *Not What You’ve Signed up for: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection*.

<https://arxiv.org/abs/2302.12173>.

MITRE Corporation. 2024. *MITRE ATLAS: Adversarial Threat Landscape for Artificial-Intelligence Systems*. <https://atlas.mitre.org/>.

OWASP Foundation. 2025. *OWASP Top 10 for Large Language Model Applications*. <https://owasp.org/www-project-top-10-for-large-language-model-applications/>.

Perez, Fábio, and Ian Ribeiro. 2022. *Ignore Previous Prompt: Attack Techniques for Language Models*. <https://arxiv.org/abs/2211.09527>.

PyCQA. 2024. *Bandit: A Security Linter for Python Source Code*.

<https://bandit.readthedocs.io/>.

Semgrep, Inc. 2024. *Semgrep: Lightweight Static Analysis for Many Languages*.

<https://semgrep.dev/>.

Willison, Simon. 2022. *Prompt Injection Attacks Against GPT-3*.

<https://simonwillison.net/2022/Sep/12/prompt-injection/>.